

# A Systems Engineering Analysis of Dual Energy Recovery in the Porsche 919 Hybrid

Ben Thomas  
CSE 185 – Spring 2026  
University of California, Santa Cruz

## Abstract

The Porsche 919 Hybrid was the only Le Mans Prototype 1 (LMP1-H) car to recover energy from both braking and the exhaust gas stream, pairing a front-axle kinetic recovery system with a generator driven by surplus exhaust flow. This paper examines the 919 from a systems engineering perspective and argues that its dominance was as much a control and software achievement as a mechanical one. We will look at the LMP1-H efficiency regulations, the combustion and hybrid powertrain, the architecture that coordinated two independent energy sources, and the lap-level energy management problem of deploying a fixed energy budget for minimum lap time. This paper assumes an embedded systems background and no motorsport knowledge; motorsport terms are defined as introduced.

## Index Terms

hybrid powertrain, energy management, optimal control, embedded systems, motorsport, energy recovery

## I. INTRODUCTION

The FIA World Endurance Championship (WEC), a worldwide championship for sports prototype and grand touring vehicles with the 24 Hours of Le Mans as its main event, confirmed Porsche's return to the premier class in 2012. From 2014 to 2017, the 919 Hybrid, the vehicle it created specifically for this purpose, participated in the Le Mans Prototype 1 Hybrid (LMP1-H) class. The 919 won Le Mans three times in a row, from 2015 to 2017, and won the WEC manufacturers' and drivers' titles in each of those seasons after placing third overall in its 2014 debut season [1], [6].

What set the 919 apart from its rivals was its powertrain. Where the competing LMP1-H cars from Audi and Toyota recovered energy only from braking, the 919 recovered energy from two independent sources: a motor-generator on the front axle that captured kinetic energy under braking, and a second generator driven by surplus flow in the car's exhaust system [1], [5]. It was the only LMP1-H car to extract energy from both the braking and the acceleration phases of a lap. After the program ended, a derestricted development car, the 919 Hybrid Evo, set a lap record of 5:19.55 at the Nürburgring Nordschleife in 2018, beating a mark that had stood since 1983 [8].

This paper examines the 919 from a systems engineering perspective. The central argument is that the car's dominance rested as much on control and software as on mechanical design. Two independent energy sources, a fixed fuel allowance, and a capped per-lap energy budget together defined a constrained optimization problem: decide how and when to deploy stored energy so that lap time is minimized and no budget is exceeded. The sections that follow describe the regulations that created this problem (Section II), the combustion and hybrid hardware that solved it physically (Sections III and IV), the architecture that coordinated the two energy sources (Section V), and the energy management strategy that governed deployment over a lap (Section VI).

## II. THE LMP1-H REGULATORY FRAMEWORK

The Automobile Club de l'Ouest (ACO) and the Federation Internationale de l'Automobile (FIA) jointly issued the LMP1-H regulations, which were based on a single efficiency principle: the more energy an automobile used from its recovery systems, the less fuel it was allowed to burn [1], [10]. Fuel consumption was restricted by a certain amount of fuel energy permitted every lap and a required fuel-flow meter that limits the rate of gasoline given to the engine, rather than by the size of the tank. Because competitiveness was directly tied to how well the engine and its hybrid systems converted energy into lap time, manufacturers were free to select nearly any engine architecture, whether diesel or gasoline, normally aspirated or turbocharged.

Hybrid systems were grouped into four energy classes, defined by the amount of recovered energy a car could deploy per lap: 2, 4, 6, or 8 megajoules (MJ) [2]. The class figure was referenced to the 13.629 km lap at Le Mans and scaled for every other circuit by track length. Selecting a higher class meant accepting a smaller fuel allowance, as illustrated in Fig. 1, so a manufacturer had to be confident that the extra deployable energy more than paid for the fuel it gave up. The 919 ran in the 6 MJ class in its 2014 debut season, then moved to the top 8 MJ class from 2015 onward, the only car ever to compete there [2].




Fig. 1. The LMP1-H efficiency rule (conceptual illustration, not drawn from a specific data source). Selecting a higher hybrid energy class lowered the fuel energy a car was allowed per lap, so adding a second recovery source became a direct competitive lever. The 919 operated in the 6 MJ class in 2014 and moved to the 8 MJ class from 2015 onward [2].

TABLE I  
LMP1-H HYBRID CONCEPTS

|                  | Porsche 919     | Audi R18      | Toyota TS040/050 |
|------------------|-----------------|---------------|------------------|
| Energy class     | 8 MJ            | 6 MJ (varied) | 6/8 MJ           |
| Recovery sources | Brake + exhaust | Brake         | Brake (2 axles)  |
| Storage          | Li-ion (800 V)  | Flywheel      | Supercapacitor   |
| Driven axles     | Rear + front    | Rear + front  | Rear + front     |

The rules required at least one energy recovery system and permitted more than one [1]. This is what made the 919's dual-recovery concept legal. A second recovery source let Porsche reach the 8 MJ class without resorting to an impractically large single recovery device. Table I compares the three manufacturers' approaches. Audi recovered braking energy alone and stored it in a flywheel; Toyota recovered from two axles into a supercapacitor; Porsche alone combined braking recovery with exhaust recovery and stored the result in a lithium-ion battery [4]. The per-lap budgets were circuit specific. At Silverstone in 2016, for instance, the 919 was allowed roughly 1.51 kg of fuel and 5.37 MJ of recovered energy per lap [3].

### III. COMBUSTION PLATFORM

A 2.0-liter turbocharged V4 engine with a 90-degree angle between its cylinder banks powered the 919 [4]. With a single Garrett turbocharger, four valves per cylinder, direct fuel injection, dry-sump lubrication, and double overhead camshafts, it reached a speed of about 9,000 rpm. The engine was a stressed element of the chassis since the aluminum crankcase was entirely load-bearing; rather than resting in a separate frame, it was attached directly between the rear suspension and the cockpit structure, carrying structural stresses on its own. Porsche claimed that the engine's 500 horsepower in racing trim was the most efficient combustion engine it had ever developed [6].

The regulations directly influenced the decision to use such a small engine. Raw displacement had little benefit since fuel flow and per-lap fuel energy were limited; what was important was getting the maximum work out of each fuel unit. Porsche was able to fulfill the fuel limit while allocating as much of the car's overall power budget to the hybrid system thanks to a downsized, compact, highly boosted engine. The overall output increased to well over 900 horsepower when the front motor released its stored energy [1].

The feature that linked the combustion engine to the second recovery system was its split exhaust. In a conventional turbocharged engine, a wastegate is a valve that bleeds off excess exhaust pressure to keep the turbocharger

from overspeeding, releasing that energy as waste heat. The 919 instead routed exhaust flow along two paths (Fig. 2): one drove the turbocharger that compressed the intake air, while the other drove a separate turbine coupled to an electrical generator [1], [5]. That generator turbine took over the role the wastegate would otherwise play, harvesting surplus exhaust energy rather than discarding it. The combustion platform was therefore not only a power source but also one of the car's two energy recovery sites, a point developed in Section IV.




Fig. 2. Split exhaust architecture of the Porsche 919, based on descriptions in [1], [5]. Exhaust gas (red) divides between the turbocharger, which compresses intake air for boost, and a separate recovery turbine that drives an electrical generator (teal) feeding the battery. The recovery turbine replaces the wastegate found in conventional turbocharged engines.

#### IV. HYBRID POWERTRAIN ARCHITECTURE

The subsystem connections of the 919 are depicted in Fig. 3. The rear axle was directly driven by the combustion engine. Under braking, the front-axle motor-generator unit (MGU) gathered kinetic energy and converted it to electrical energy stored in the battery. On demand, the MGU reversed its role and used that stored energy to drive the front wheels, momentarily putting the vehicle in all-wheel drive on corner exits. The same battery was replenished by the exhaust generator turbine described in Section III. A central, liquid-cooled lithium-ion battery running at 800 V was charged by both recovery routes [4]. Porsche stated that roughly 40 percent of recovered energy came from the exhaust generator turbine and about 60 percent from braking [1].




Fig. 3. Porsche 919 Hybrid powertrain architecture, reconstructed from published descriptions [1], [4]. The combustion engine drives the rear axle; a front-axle motor-generator (MGU) recovers braking energy and deploys it to the front axle for temporary all-wheel drive. A separate exhaust generator turbine harvests surplus exhaust energy. Both electrical sources feed a central 800 V lithium-ion battery.

The 800 V battery voltage was a deliberate engineering choice. At a given power level, doubling voltage halves current, which reduces resistive losses in wiring and power electronics and permits lighter cabling. In 2014, 800 V was virtually unheard of in either motorsport or production vehicles; the industry standard for battery electric road cars at the time was roughly 400 V. Porsche later stated that the 919 program directly informed the 800 V architecture of its Mission E concept car, which became the Taycan production vehicle [1], [9]. The 919 thus served as a technology demonstrator whose electrical architecture reached series production.

For Porsche, the front-axle kinetic energy recovery system (KERS) was not a novel concept. A flywheel-based KERS was previously tried on the 911 GT3 R Hybrid, a GT-class race vehicle that participated in a few endurance races between 2010 and 2011 [5]. Instead of using a battery for storage, that design employed a mechanical flywheel. Porsche increased energy density and packaging flexibility for the 919 by swapping out the flywheel for a lithium-ion battery and the back axle for a front axle motor. In addition to the engine's rear-axle output, the front-axle motor could provide about 400 horsepower during deployment [1].

The battery was situated in the middle of the monocoque, between the engine and the cockpit. Cell temperatures were kept within a limited operational window by liquid cooling. The battery had to meet two conflicting requirements: it had to be light enough to avoid penalizing the car's minimal weight and have enough capacity to hold the whole deployment budget for a single lap. Due to these limitations, the high energy density of lithium-ion chemistry and the 800 V working voltage were crucial to the design.

## V. CONTROL ARCHITECTURE

The ECU hardware and control software of the production controller, which managed the 919's two energy sources, were never made public by Porsche. As a result, we use two publicly available lenses to study the control problem: the driver-facing interface and the energy management controller structure found in the academic literature. The two-level structure typical of such systems is depicted in Fig. 4, where an offline supervisory layer computes the deployment plan and an online execution layer carries it out in real time.




Fig. 4. Generic two-level control structure for hybrid race powertrains, as described in the academic literature [11], [12], [14]. This is not a diagram of the 919's proprietary controller, which Porsche never published. An offline supervisory layer computes the energy-deployment strategy for the lap; an online layer executes it and corrects against measured battery state of charge (SoC), fuel flow, and vehicle speed.

### A. Driver Interface

In essence, the 919's steering wheel was a portable computer [7]. It had rotary switches for selecting traction control sensitivity, hybrid deployment programs, and engine maps (which controlled fuelling, ignition, and boost objectives). The driver may manually start or stop recovery and deployment actions by using paddles behind the wheel. In actuality, the ECU made cycle-by-cycle judgments while the driver chose high-level settings. This separation of duties is similar to the supervisory/execution divide seen in Fig. 4: the online controller converted the strategic parameters provided by the driver via the steering wheel into actuator orders at millisecond intervals.

### B. Brake-by-Wire and Torque Blending

The brake-by-wire system was one of the most important subsystems for safety. The hydraulic friction brakes and the front MGU's regenerative torque both applied torque to the front axle at the same time while braking. No matter how much energy the MGU was recovering at any given time, the ECU had to combine these two sources to provide the driver with a constant brake pedal sensation. The MGU contribution decreased to zero and the friction brakes had to instantaneously adapt if the battery was full and could take no more charges. This blending problem is a real-time control task with a strict deadline; the driver perceives any delay or mismatch as an inconsistent braking force, which is a major safety risk at 300 km/h.

### C. Rear-Wheel Drive and All-Wheel Drive Transitions

During each lap, the 919 alternated between all-wheel drive (AWD) and rear-wheel drive (RWD). The car became AWD when the front MGU transferred stored energy to the front wheels during straight-line acceleration. The car ran in RWD on engine power alone during turns, with the front motor usually inactive. The ECU used the remaining energy budget, the battery's current state of charge (SoC), and the traction at the front axle to determine whether to engage the front motor when the driver pushed the throttle during a corner departure. Traction control, also known as anti-slip regulation, or ASR, kept an eye on the speeds of the front and rear wheels and stopped deployment if the front wheels started to spin. This entire decision loop, from traction sensing to motor torque adjustment, had to execute within a few milliseconds to prevent wheel spin from destabilizing the car.

### D. Telemetry and the Data Pipeline

The control architecture depended on high-bandwidth telemetry. Sensors throughout the car measured battery voltage, cell temperatures, motor currents, wheel speeds, fuel flow, and dozens of other parameters. This data was logged on board and transmitted to the pit wall during sessions. Sohn et al. describe a general framework for organizing the collection and use of such vehicle data in automotive systems engineering, structuring it by purpose (the “Why”), operational context (the “Where”), development phase (the “When”), data sources (the “What”), and processing method (the “How”) [15]. In a motorsport context, the cycle is compressed: telemetry from a practice session feeds ECU calibration changes that are deployed for qualifying the same day. The validity and coverage of the data increase as the car moves from simulation to physical prototype to on-track testing, which maps directly onto the V-model development progression that Sohn et al. describe [15]. The 919's control calibration was therefore not a static artifact but a continuously refined product of a data-driven engineering loop.

## VI. ENERGY MANAGEMENT STRATEGY

The regulations turned racing into a resource-allocation problem. Once the hardware is fixed, the per-lap fuel allowance sets roughly how much work the engine can do, and the energy class sets how much electrical energy the car may deploy on top of it. Neither budget can be exceeded, and the battery must end each lap at the charge it began with so the same plan repeats lap after lap. The freedom that remains is *where* on the lap to spend the electrical energy. Deploying electric power adds to the engine's output at the wheels, which raises speed and lowers lap time, so the strategy is to place a fixed budget where it buys the most time.

### A. Problem Formulation

The decision variable is the deployment power profile  $P_{\text{dep}}(t)$  over one lap. The objective is to minimize lap time:

$$\min_{P_{\text{dep}}(\cdot)} T_{\text{lap}}(P_{\text{dep}}). \quad (1)$$

This minimization is subject to five constraints. The first two are budget constraints. The total fuel energy consumed over the lap must not exceed the regulatory allowance  $E_{\text{fuel}}^{\text{max}}$ , and the total electrical energy deployed must not exceed the class limit  $E_{\text{hyb}}^{\text{max}}$ :

$$\int_{\text{lap}} P_{\text{fuel}}(t) dt \leq E_{\text{fuel}}^{\text{max}}, \quad (2)$$

$$\int_{\text{lap}} P_{\text{dep}}(t) dt \leq E_{\text{hyb}}^{\text{max}}. \quad (3)$$

The third constraint describes the battery dynamics. The stored energy  $E(t)$  rises when recovery power  $P_{\text{rec}}$  flows in and falls when deployment power  $P_{\text{dep}}$  flows out:

$$E(t) = P_{\text{rec}}(t) - P_{\text{dep}}(t). \quad (4)$$

The fourth constraint keeps the battery within its physical limits at all times, where  $E_{\text{batt}}$  is the total pack capacity:

$$0 \leq E(t) \leq E_{\text{batt}}. \quad (5)$$

The fifth constraint requires the battery to end each lap at the same charge it started with, so the same strategy can repeat every lap throughout a race stint:

$$E(T_{\text{lap}}) = E(0). \quad (6)$$

The fact that this formulation differs from the energy management problem in a passenger automobile is one of its key characteristics. A road vehicle's hybrid system aims to consume less fuel while maintaining the speed that the driver or driving cycle determines. The velocity profile is an input rather than a decision variable. When a racing car uses electric boost on a straight, it raises its physical speed, which changes the recovery opportunity when it enters the next braking zone [11]. The velocity profile and the energy strategy must be adjusted simultaneously since they are connected.

### B. Convex Reformulation and Solution

Ebbesen et al. address this problem for a Formula 1 hybrid power unit, which shares the same structural elements as the 919: a turbocharged combustion engine, a kinetic MGU on the drivetrain, a second generator on the turbocharger (analogous to the 919's exhaust generator turbine), and a lithium-ion battery [11]. A separation argument is the central idea of their methodology. The driver uses steering inputs to select the racing line, which is the geometric route the car takes around the circuit. It is unchangeable by the energy management system. Instead of jointly optimizing the path, this allows the optimizer to solve for the velocity profile and the power split along a given, defined path. The issue size is significantly reduced by such separation.

Ebbesen et al. reformulate the issue as a second-order conic program (SOCP) by adding convex approximations and relaxations to the vehicle dynamics. Convexity ensures that whatever solution the solver discovers is the global optimum rather than just a local one that may be enhanced by beginning from a different place. Additionally, the computation is quick: on a 2.6 GHz CPU, a whole lap optimization takes less than a second to complete. In a total of around six minutes, the authors covered 357 permutations of fuel load and battery consumption target [11]. During a race weekend, when engineers have little time between sessions to modify the car's strategy, this type quick turnaround is quite helpful. Sohn et al. would classify these parameter sweeps as prescriptive analytics within their "How" layer: model outputs that determine what the ECU calibration should be [15].

### C. Optimal Deployment Policy

The strategic findings from Ebbesen et al. are clear and apply directly to the 919. First, the optimal use of the kinetic MGU is bang-bang: deploy full electric boost at the start of each straight where the car is slow and the speed gain per unit of energy is largest, then shut the motor off. Attempting to spread the boost more evenly across the lap is slower [11]. Second, over the course of multiple laps, the best approach is to keep the battery roughly charge neutral, recovering about as much energy each lap as is spent. Charging the battery extra, by braking harder or recovering over longer braking zones, costs more lap time than the additional boost energy is worth. This is because the conversion losses in recovery and redeployment are non-negligible, and the aerodynamic drag penalty for carrying extra speed into a braking zone grows quadratically with velocity [11].

These results suggest that the controller for the 919 should activate the entire front-axle motor power right after each corner exit, maintain it until the energy slice for that straight is depleted, and then switch back to engine-only rear-wheel drive for the duration of the straight. Automatic recovery takes place in the next braking zone. The resulting proposed battery SoC trend over a single lap is depicted in Fig. 5: charge builds during braking and declines during deployment zones before recovering to its initial value at the end of the loop.




Fig. 5. Conceptual battery state of charge (SoC) trajectory over one lap (illustrative, not measured 919 data). Charge falls during deployment zones (straights) and rises during braking recovery. The charge-sustaining constraint requires SoC to return to its starting value at lap end. The sawtooth pattern reflects the bang-bang deployment policy described by Ebbesen et al. [11].

#### D. Online Controller

The previously mentioned offline technique generates a reference deployment plan that specifies how much energy should be recovered in each braking zone and how much should be deployed on each straight. This schedule is carried out by the online controller, which operates on the ECU at the millisecond level and adjusts for real-time disruptions. The SoC at the beginning of the subsequent straight will be different from the plan if the driver stops a little earlier than anticipated. The deployment period must be modified appropriately by the online layer. The controller determines the intended deployment for the current segment, reads the current SoC and track location, and modifies power output according to whether the battery is ahead of or behind the reference trajectory. This reasoning is depicted in Algorithm 1. This structure corresponds to a receding-horizon or model predictive control (MPC) approach, as described by Salazar et al. [14]. The pseudocode is a simplified conceptual representation of the general control approach described in the literature, not a depiction of Porsche’s proprietary implementation.

**Algorithm 1** Online energy management controller (conceptual, based on the two-level MPC structure in [14]). The controller corrects the offline deployment schedule against measured SoC at each control step.

---

```

1: for each control step ( $\Delta t \approx 1$  ms) do
2:   Read  $\text{SoC}_{\text{actual}}$ , position, speed
3:   segment  $\leftarrow$  LOOKUPSEGMENT(position)
4:    $\text{SoC}_{\text{ref}} \leftarrow$  OFFLINEPLAN(segment)
5:   error  $\leftarrow \text{SoC}_{\text{actual}} - \text{SoC}_{\text{ref}}$ 
6:   if segment is a straight then
7:     if  $\text{SoC}_{\text{actual}} > \text{SoC}_{\text{min}}$  then
8:        $P_{\text{deploy}} \leftarrow P_{\text{max, MGU}}$ 
9:       if error  $< -\text{threshold}$  then
10:         $P_{\text{deploy}} \leftarrow P_{\text{deploy}} \times (1 - k \cdot |\text{error}|)$ 
11:      end if
12:    else
13:       $P_{\text{deploy}} \leftarrow 0$ 
14:    end if
15:  else if segment is a braking zone then
16:     $P_{\text{recover}} \leftarrow \min(P_{\text{max, regen}}, F_{\text{brake}} \times r_{\text{blend}})$ 
17:     $P_{\text{friction}} \leftarrow F_{\text{brake}} - P_{\text{recover}}$ 
18:  end if
19:  Apply( $P_{\text{deploy}}, P_{\text{recover}}, P_{\text{friction}}$ )
20: end for

```

---

#### E. Worked Example: Budget Scaling

The 8 MJ deployment cap referenced the 13.629 km Le Mans lap and scaled linearly by track length for other circuits. At a shorter circuit such as the 6.004 km Nürburgring sprint layout used in the 2015 WEC round, the

deployable energy budget would scale to approximately  $8 \times (6.004/13.629) \approx 3.52$  MJ per lap. The fuel allowance would scale similarly. This means the controller had to re-solve or re-tune the deployment schedule for every circuit on the calendar. The fast solve times enabled by convex formulations, as demonstrated by Ebbesen et al., make this re-tuning practical even during a compressed race-weekend schedule [11]. Duhr et al. extend this analysis to show that energy management also interacts with gearshift strategy: the choice of which gear to use on corner exit affects how much energy can be recovered during the subsequent braking event, adding another coupled decision variable [13].

## VII. CONCLUSION

The dual-recovery concept was viable because the regulations rewarded it directly: a second recovery site let the 919 reach the 8 MJ class and trade fuel for hybrid energy more aggressively than any rival. But the hardware alone did not produce the results. The control architecture that coordinated two independent energy sources, blended regenerative and friction braking in real time, and managed RWD-to-AWD transitions multiple times per lap was equally essential. The energy management strategy on top of it, placing a fixed energy budget where it bought the most lap time, was a constrained optimization problem whose solution required careful offline computation and robust online execution.

The broader lesson for hybrid systems engineering is one of co-design. The regulations, the hardware, and the control strategy were developed together. No single element in isolation would have been sufficient. A dual-recovery powertrain without an effective deployment strategy would leave energy on the table. A sophisticated controller without a second recovery source would not have reached the 8 MJ class. The offline-strategy, online-execution pattern visible in Fig. 4 recurs throughout embedded control and is not unique to motorsport. What is distinctive about the 919's case is the severity of the constraints: fixed energy budgets, charge-sustaining requirements, and lap times measured in fractions of a second combined to make the control problem as demanding as the mechanical design.

The data infrastructure that supported this control loop deserves mention as well. Sohn et al. argue that competitive advantage in modern automotive engineering lies at the intersection of validated control software, systematic telemetry analysis, and strategies that approach the theoretical optimum [15]. The 919 program exemplified this intersection. Telemetry from practice sessions fed parameter sweeps that updated the deployment strategy; the strategy was loaded onto the ECU and validated on track; and the results fed back into the next iteration. This closed loop, running on a compressed race-weekend timeline, is a concrete instance of the data-driven engineering process that Sohn et al. formalize.

Some of the 919's technology reached production. Porsche has stated that the car's 800 V architecture fed into the Mission E concept and the Taycan [1], [9]. The broader motorsport landscape also evolved. Table II contrasts the open LMP1-H rules with the cost-capped Hypercar regulations (Le Mans Hypercar, or LMH, and Le Mans Daytona hybrid, or LMDh) that replaced them starting in 2021.

TABLE II  
FROM LMP1-H TO THE HYPERCAR ERA

|                  | <b>LMP1-H (919 era)</b>        | <b>LMH / LMDh</b>            |
|------------------|--------------------------------|------------------------------|
| Hybrid freedom   | Open architecture              | Constrained / spec hybrid    |
| Recovery sources | Unrestricted (brake + exhaust) | Limited (rear only for LMDh) |
| Hybrid drive     | Front axle (AWD)               | Rear (LMDh) / either (LMH)   |
| Cost philosophy  | Unbounded R&D                  | Cost cap / BoP equalization  |
| Energy class     | Up to 8 MJ                     | Lower, fixed targets         |

The Hypercar rules deliberately constrain the engineering freedom that the 919 exploited. LMDh cars use a spec hybrid unit supplied by a single manufacturer, with recovery limited to the rear axle and deployment capped well below 8 MJ. LMH cars have slightly more freedom but operate under a Balance of Performance (BoP) system that adjusts power and weight to equalize competitiveness across manufacturers. Neither format permits the kind of unbounded, dual-recovery optimization that defined the 919 program. The regulations shifted from rewarding engineering depth to controlling costs and ensuring competitive parity, which reduced the role of control software as a differentiator.



The 919 Hybrid thus occupies a specific historical moment: the last era in which endurance racing regulations explicitly rewarded engineering complexity in hybrid powertrain design. The car's three consecutive Le Mans victories from 2015 to 2017 were the product of a tightly integrated system in which combustion, electrical, and control engineering reinforced one another. For embedded systems engineers, the 919 remains a compelling case study in how software and control strategy can be as decisive as mechanical hardware.

## REFERENCES

**Note: The Porsche Newsroom sources are short articles and technical documents, which is the reason for having quite a few of them**

- [1] Porsche AG, "How the technology of the 919 Hybrid works," Porsche Newsroom, Jul. 2016. [Online]. Available: <https://newsroom.porsche.com/en/motorsports/porsche-motorsport-fiawec-919-hybrid-technology-lmp1-race-car-12724.html>
- [2] Porsche AG, "Porsche heads to Silverstone full of energy," Porsche Newsroom, 2015. [Online]. Available: <https://newsroom.porsche.com/en/motorsports/porsche-silverstone-wec-919-hybrid-2015-le-mans-10159.html>
- [3] Porsche AG, "Porsche 919 Hybrid benefits from efficiency," Porsche Newsroom, Apr. 2016. [Online]. Available: <https://newsroom.porsche.com/en/motorsports/porsche-motorsport-silverstone-911-rsr-911-gt3r-wec-12424.html>
- [4] Porsche AG, "Technical check: 919 Hybrid Evo," Porsche Newsroom (Top 5 series), Jan. 2019. [Online]. Available: <https://newsroom.porsche.com/en/motorsports/porsche-919-hybrid-evo-top-5-series-technical-check-16834.html>
- [5] Porsche AG, "Powered by two hearts," Porsche Newsroom, Oct. 2017. [Online]. Available: <https://newsroom.porsche.com/en/motorsports/porsche-powered-by-two-hearts-performance-hybrid-vehicles-13618.html>
- [6] Porsche AG, "World champion turbo four-cylinder," Porsche Newsroom, Feb. 2016. [Online]. Available: <https://newsroom.porsche.com/en/innovation/engineering/porsche-motorsports-turbo-four-cylinder-engine-919-hybrid-12226.html>
- [7] Porsche AG, "Multifunctional control centre," Porsche Newsroom. [Online]. Available: <https://newsroom.porsche.com/en/innovation/engineering/multifunctional-control-centre-10727.html>
- [8] Porsche AG, "5:19.55 minutes – Porsche 919 Hybrid Evo takes record in the Green Hell," Porsche Newsroom, Jun. 2018. [Online]. Available: <https://newsroom.porsche.com/en/motorsports/porsche-919-hybrid-evo-record-nuerburgring-nordschleife-5-minutes-19-seconds-55-timo-bernhard-15752.html>
- [9] Porsche AG, "LMDh regulations: Dawn of a new era in long-distance motorsport," Porsche Newsroom, 2023. [Online]. Available: <https://newsroom.porsche.com/en/press-kits/le-mans-2023/LMDh-regulations-Dawn-of-a-new-era-in-long-distance-motorsport.html>
- [10] Fédération Internationale de l'Automobile and Automobile Club de l'Ouest, "LMP Technical Regulations," FIA/ACO. [Online]. Available: <https://www.fia.com/file/75500/download/24356?token=Dr4gK1cJ>
- [11] S. Ebbesen, M. Salazar, P. Elbert, C. Bussi, and C. H. Onder, "Time-optimal control strategies for a hybrid electric race car," *IEEE Trans. Control Syst. Technol.*, vol. 26, no. 1, pp. 233–247, Jan. 2018, doi: 10.1109/TCST.2017.2661824.
- [12] M. Salazar, P. Elbert, S. Ebbesen, C. Bussi, and C. H. Onder, "Time-optimal control policy for a hybrid electric race car," *IEEE Trans. Control Syst. Technol.*, vol. 25, no. 6, pp. 1921–1934, 2017, doi: 10.1109/TCST.2016.2642830.
- [13] P. Duhr, G. Christodoulou, C. Balerna, M. Salazar, A. Cerofolini, and C. H. Onder, "Time-optimal gearshift and energy management strategies for a hybrid electric race car," *Applied Energy*, vol. 282, art. 115980, 2021, doi: 10.1016/j.apenergy.2020.115980.
- [14] M. Salazar, C. Balerna, P. Elbert, F. P. Grandi, and C. H. Onder, "Real-time control algorithms for a hybrid electric race car using a two-level model predictive control scheme," in *Proc. IEEE Conf. Decision Control (CDC)*, Melbourne, Australia, Dec. 2017, pp. 3574–3581, doi: 10.1109/CDC.2017.8264185.
- [15] T. S. Sohn, M. Dillitzer, T. Brühl, R. Schwager, T. D. Eberhardt, P. Elspas, and E. Sax, "The why, where, when, what and how of data for data-driven engineering of automotive systems," in *Proc. IEEE Int. Symp. Syst. Eng. (ISSE)*, 2024, pp. 1–8, doi: 10.1109/ISSE63315.2024.10741122.